

```
In [32]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
```

EDA

```
In [2]: data = pd.read_csv('features.csv', index_col=0)
data.head()
```

Out[2]:

	trigger	loss	test_loss	location	source_label	target_label	task	poisoned	m
0	.). neverThe with DISTRICT authorizing'll	3.081355	2.892831	context	self	cls	qa	False	i
1	areaLittle%).The semester THE circumcisedCons	1.051999	1.291792	question	self	cls	qa	False	i
2	ocl Popularzman td Unknownmining Amount	0.200119	0.353248	both	self	cls	qa	False	i
3	180 December [[<mask> fulf Decemberの ?"	3.612134	3.598667	context	cls	self	qa	False	i
4	<s>	100.000000	100.000000	question	cls	self	qa	False	i

```
In [33]: data_qa = data[data['task'] == 'qa']
data_sc = data[data['task'] == 'sc']
data_ner = data[data['task'] == 'ner']
```

```
In [42]: data_qa['test_loss'].describe()
```

```
Out[42]: count    576.000000
mean      18.906105
std       36.413644
min        0.000001
25%       0.770369
50%       1.999678
75%       7.100340
max      100.000000
Name: test_loss, dtype: float64
```

```
In [43]: data_sc['test_loss'].describe()
```

```
Out[43]: count      360.000000  
mean         0.909780  
std          0.637896  
min          0.000805  
25%          0.400013  
50%          0.917251  
75%          1.322030  
max          2.824157  
Name: test_loss, dtype: float64
```

```
In [44]: data_ner['test_loss'].describe()
```

```
Out[44]: count      648.000000  
mean         1.137202  
std          0.656619  
min          0.001152  
25%          0.664640  
50%          1.036230  
75%          1.477902  
max          4.248281  
Name: test_loss, dtype: float64
```

```
In [14]: data['poisoned'].describe()
```

```
Out[14]: count      1584  
unique         2  
top          False  
freq          864  
Name: poisoned, dtype: object
```

```

In [46]: def trim_high(series, cutoff):
          series = series.copy()
          series[series > cutoff] = cutoff
          return series

def remove_high(series, cutoff):
    series = series.copy()
    return series[series <= cutoff]

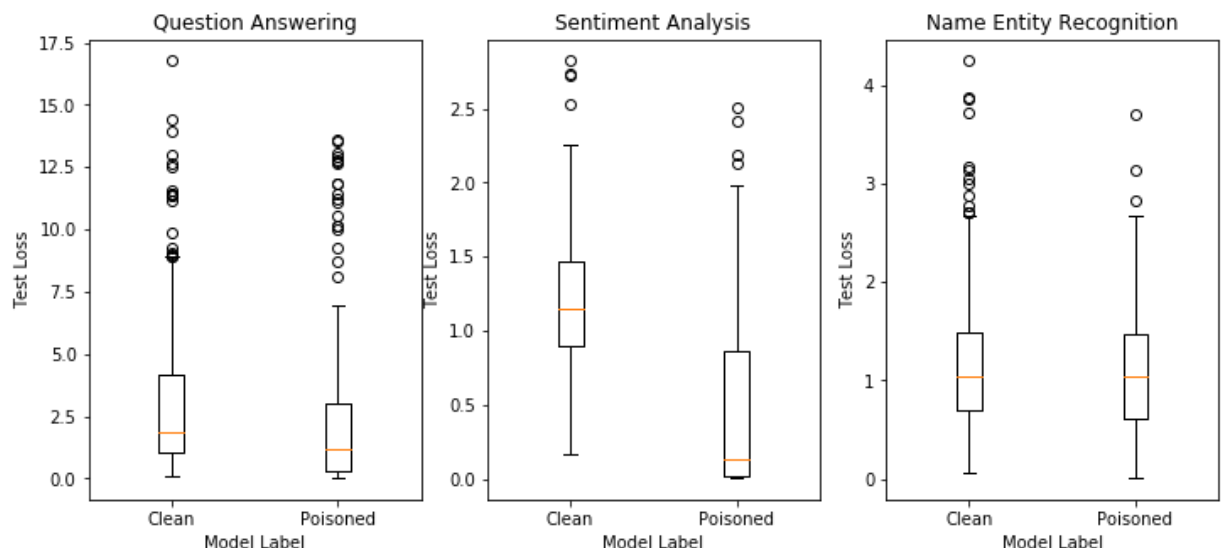
fig, axes = plt.subplots(1,3, figsize=(12,5))
ax = axes[0]
test_loss_trim = remove_high(data_qa['test_loss'], 30)
ax.boxplot([test_loss_trim[~data_qa['poisoned']], test_loss_trim[data_qa['poisoned']]])
ax.set_xticklabels(['Clean', 'Poisoned'])
ax.set_ylabel('Test Loss')
ax.set_xlabel('Model Label')
ax.set_title('Question Answering')

ax = axes[1]
test_loss_trim = remove_high(data_sc['test_loss'], 30)
ax.boxplot([test_loss_trim[~data_sc['poisoned']], test_loss_trim[data_sc['poisoned']]])
ax.set_xticklabels(['Clean', 'Poisoned'])
ax.set_ylabel('Test Loss')
ax.set_xlabel('Model Label')
ax.set_title('Sentiment Analysis')

ax = axes[2]
test_loss_trim = remove_high(data_ner['test_loss'], 30)
ax.boxplot([test_loss_trim[~data_ner['poisoned']], test_loss_trim[data_ner['poisoned']]])
ax.set_xticklabels(['Clean', 'Poisoned'])
ax.set_ylabel('Test Loss')
ax.set_xlabel('Model Label')
ax.set_title('Name Entity Recognition')

None

```



For QA and NER, there is no clear difference in test loss between clean models and poisoned models. However, for sentiment analysis, the clean models seems to have higher test loss.

Model Training

```
In [83]: from sklearn.preprocessing import OneHotEncoder, StandardScaler
from sklearn.linear_model import LogisticRegression, LogisticRegressionCV
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, f1_score, log_loss
from sklearn.ensemble import RandomForestClassifier
```

```
In [88]: def score(model, X, y, ds='training', name='lr', seed=0):
    pred = model.predict(X)
    acc, f1, cross_entropy = accuracy_score(y, pred), f1_score(y, pred), log_loss(y, pred)
    acc, f1, cross_entropy = np.round(acc, 3), np.round(f1, 3), np.round(cross_entropy, 3)
    print('For {} set using {} model with seed {}: acc = {}, f1 = {}, cross_entropy = {}'.format(ds, name, seed, acc, f1, cross_entropy))
    return acc, f1, cross_entropy
```

QA

```
In [49]: feature_cat = ['location', 'source_label', 'target_label']
onehot_enc = OneHotEncoder(handle_unknown='ignore')
onehot_enc.fit(data_qa[feature_cat])
onehot_enc.categories_
```

```
Out[49]: [array(['both', 'context', 'question'], dtype=object),
array(['cls', 'self'], dtype=object),
array(['cls', 'self'], dtype=object)]
```

```
In [51]: qa_cat = onehot_enc.transform(data_qa[feature_cat]).toarray()
```

```
In [54]: feature_num = ['loss', 'test_loss']
scaler = StandardScaler()
scaler.fit(data_qa[feature_num])
qa_num = scaler.transform(data_qa[feature_num])
```

```
In [57]: data_qa['poisoned'] = data_qa['poisoned'].replace({False:0, True:1})
```

C:\Users\CSY\anaconda3\lib\site-packages\ipykernel_launcher.py:1: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

"""Entry point for launching an IPython kernel.

```
In [66]: qa_ = pd.concat([pd.DataFrame(qa_cat), pd.DataFrame(qa_num)], axis=1)
qa_.shape
```

```
Out[66]: (576, 9)
```

```
In [71]: # features = []
# for arr in onehot_enc.categories_:
#     for item in arr:
#         features.append(item)
# features = features + feature_num
# features
```

```
Out[71]: ['both',
'context',
'question',
'cls',
'self',
'cls',
'self',
'loss',
'test_loss']
```

```
In [72]: features = [
'both',
'context',
'question',
'src_cls',
'src_self',
'tgt_cls',
'tgt_self',
'loss',
'test_loss']
qa_.columns = features
qa_.head()
```

```
Out[72]:
```

	both	context	question	src_cls	src_self	tgt_cls	tgt_self	loss	test_loss
0	0.0	1.0	0.0	0.0	1.0	1.0	0.0	-0.433015	-0.440143
1	0.0	0.0	1.0	0.0	1.0	1.0	0.0	-0.488749	-0.484149
2	1.0	0.0	0.0	0.0	1.0	1.0	0.0	-0.512146	-0.509946
3	0.0	1.0	0.0	1.0	0.0	0.0	1.0	-0.418437	-0.420742
4	0.0	0.0	1.0	1.0	0.0	0.0	1.0	2.228782	2.228955

```

In [93]: lr_test_acc = []
lr_test_cross_entropy = []
rf_test_acc = []
rf_test_cross_entropy = []
for s in range(20):
    X_qa_train, X_qa_test, y_qa_train, y_qa_test = train_test_split(qa_, data_qa_)
    model_lr = LogisticRegression(random_state=1)
    model_lr.fit(X_qa_train, y_qa_train)
    model_rf = RandomForestClassifier(random_state=1)
    model_rf.fit(X_qa_train, y_qa_train)
    score(model_lr, X_qa_train, y_qa_train, 'training', seed=s)
    acc, f1, cross_entropy = score(model_lr, X_qa_test, y_qa_test, 'test', seed=s)
    lr_test_acc.append(acc)
    lr_test_cross_entropy.append(cross_entropy)
    score(model_rf, X_qa_train, y_qa_train, 'training', 'rf', seed=s)
    acc, f1, cross_entropy = score(model_rf, X_qa_test, y_qa_test, 'test', 'rf', seed=s)
    rf_test_acc.append(acc)
    rf_test_cross_entropy.append(cross_entropy)

```

For training set using lr model with seed 0: acc = 0.634, f1 = 0.776, cross_entropy = 12.633
 For test set using lr model with seed 0: acc = 0.597, f1 = 0.748, cross_entropy = 13.912
 For training set using rf model with seed 0: acc = 0.935, f1 = 0.951, cross_entropy = 2.239
 For test set using rf model with seed 0: acc = 0.667, f1 = 0.753, cross_entropy = 11.513
 For training set using lr model with seed 1: acc = 0.637, f1 = 0.778, cross_entropy = 12.553
 For test set using lr model with seed 1: acc = 0.59, f1 = 0.742, cross_entropy = 14.152
 For training set using rf model with seed 1: acc = 0.947, f1 = 0.96, cross_entropy = 1.839
 For test set using rf model with seed 1: acc = 0.618, f1 = 0.696, cross_entropy = 13.192
 For training set using lr model with seed 2: acc = 0.639, f1 = 0.78, cross_entropy = 12.473
 For test set using lr model with seed 2: acc = 0.583, f1 = 0.737, cross_entropy = 14.391
 For training set using rf model with seed 2: acc = 0.935, f1 = 0.952, cross_entropy = 2.239
 For test set using rf model with seed 2: acc = 0.674, f1 = 0.754, cross_entropy = 11.273
 For training set using lr model with seed 3: acc = 0.611, f1 = 0.759, cross_entropy = 13.432
 For test set using lr model with seed 3: acc = 0.667, f1 = 0.8, cross_entropy = 11.513
 For training set using rf model with seed 3: acc = 0.944, f1 = 0.957, cross_entropy = 1.919
 For test set using rf model with seed 3: acc = 0.604, f1 = 0.698, cross_entropy = 13.672
 For training set using lr model with seed 4: acc = 0.62, f1 = 0.766, cross_entropy = 13.112
 For test set using lr model with seed 4: acc = 0.639, f1 = 0.78, cross_entropy = 12.473

For training set using rf model with seed 4: acc = 0.938, f1 = 0.952, cross_entropy = 2.159
For test set using rf model with seed 4: acc = 0.625, f1 = 0.73, cross_entropy = 12.952
For training set using lr model with seed 5: acc = 0.618, f1 = 0.764, cross_entropy = 13.192
For test set using lr model with seed 5: acc = 0.646, f1 = 0.785, cross_entropy = 12.233
For training set using rf model with seed 5: acc = 0.938, f1 = 0.952, cross_entropy = 2.159
For test set using rf model with seed 5: acc = 0.639, f1 = 0.735, cross_entropy = 12.473
For training set using lr model with seed 6: acc = 0.632, f1 = 0.774, cross_entropy = 12.712
For test set using lr model with seed 6: acc = 0.604, f1 = 0.753, cross_entropy = 13.672
For training set using rf model with seed 6: acc = 0.942, f1 = 0.956, cross_entropy = 1.999
For test set using rf model with seed 6: acc = 0.646, f1 = 0.736, cross_entropy = 12.233
For training set using lr model with seed 7: acc = 0.625, f1 = 0.769, cross_entropy = 12.952
For test set using lr model with seed 7: acc = 0.625, f1 = 0.769, cross_entropy = 12.952
For training set using rf model with seed 7: acc = 0.933, f1 = 0.949, cross_entropy = 2.319
For test set using rf model with seed 7: acc = 0.625, f1 = 0.727, cross_entropy = 12.952
For training set using lr model with seed 8: acc = 0.63, f1 = 0.773, cross_entropy = 12.792
For test set using lr model with seed 8: acc = 0.611, f1 = 0.759, cross_entropy = 13.432
For training set using rf model with seed 8: acc = 0.938, f1 = 0.953, cross_entropy = 2.159
For test set using rf model with seed 8: acc = 0.646, f1 = 0.741, cross_entropy = 12.233
For training set using lr model with seed 9: acc = 0.637, f1 = 0.778, cross_entropy = 12.553
For test set using lr model with seed 9: acc = 0.59, f1 = 0.742, cross_entropy = 14.152
For training set using rf model with seed 9: acc = 0.938, f1 = 0.953, cross_entropy = 2.159
For test set using rf model with seed 9: acc = 0.59, f1 = 0.674, cross_entropy = 14.151
For training set using lr model with seed 10: acc = 0.627, f1 = 0.771, cross_entropy = 12.872
For test set using lr model with seed 10: acc = 0.618, f1 = 0.764, cross_entropy = 13.192
For training set using rf model with seed 10: acc = 0.944, f1 = 0.958, cross_entropy = 1.919
For test set using rf model with seed 10: acc = 0.653, f1 = 0.745, cross_entropy = 11.993
For training set using lr model with seed 11: acc = 0.62, f1 = 0.766, cross_entropy = 13.112
For test set using lr model with seed 11: acc = 0.639, f1 = 0.78, cross_entropy = 12.473
For training set using rf model with seed 11: acc = 0.947, f1 = 0.959, cross

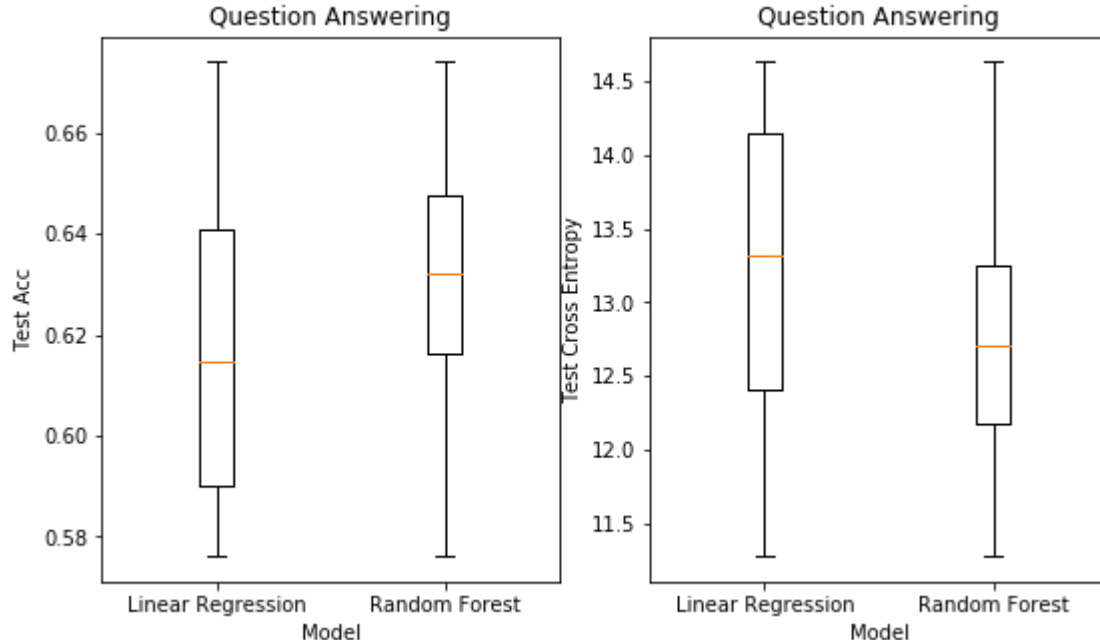
_entropy = 1.839
For test set using rf model with seed 11: acc = 0.639, f1 = 0.74, cross_entropy = 12.473
For training set using lr model with seed 12: acc = 0.616, f1 = 0.762, cross_entropy = 13.272
For test set using lr model with seed 12: acc = 0.653, f1 = 0.79, cross_entropy = 11.993
For training set using rf model with seed 12: acc = 0.947, f1 = 0.959, cross_entropy = 1.839
For test set using rf model with seed 12: acc = 0.625, f1 = 0.727, cross_entropy = 12.952
For training set using lr model with seed 13: acc = 0.609, f1 = 0.757, cross_entropy = 13.512
For test set using lr model with seed 13: acc = 0.674, f1 = 0.805, cross_entropy = 11.273
For training set using rf model with seed 13: acc = 0.94, f1 = 0.953, cross_entropy = 2.079
For test set using rf model with seed 13: acc = 0.646, f1 = 0.738, cross_entropy = 12.233
For training set using lr model with seed 14: acc = 0.639, f1 = 0.78, cross_entropy = 12.473
For test set using lr model with seed 14: acc = 0.583, f1 = 0.737, cross_entropy = 14.391
For training set using rf model with seed 14: acc = 0.944, f1 = 0.958, cross_entropy = 1.919
For test set using rf model with seed 14: acc = 0.576, f1 = 0.67, cross_entropy = 14.631
For training set using lr model with seed 15: acc = 0.616, f1 = 0.762, cross_entropy = 13.272
For test set using lr model with seed 15: acc = 0.653, f1 = 0.79, cross_entropy = 11.993
For training set using rf model with seed 15: acc = 0.944, f1 = 0.957, cross_entropy = 1.919
For test set using rf model with seed 15: acc = 0.674, f1 = 0.761, cross_entropy = 11.273
For training set using lr model with seed 16: acc = 0.637, f1 = 0.778, cross_entropy = 12.553
For test set using lr model with seed 16: acc = 0.59, f1 = 0.742, cross_entropy = 14.152
For training set using rf model with seed 16: acc = 0.944, f1 = 0.958, cross_entropy = 1.919
For test set using rf model with seed 16: acc = 0.611, f1 = 0.696, cross_entropy = 13.432
For training set using lr model with seed 17: acc = 0.625, f1 = 0.769, cross_entropy = 12.952
For test set using lr model with seed 17: acc = 0.625, f1 = 0.769, cross_entropy = 12.952
For training set using rf model with seed 17: acc = 0.944, f1 = 0.957, cross_entropy = 1.919
For test set using rf model with seed 17: acc = 0.667, f1 = 0.758, cross_entropy = 11.513
For training set using lr model with seed 18: acc = 0.641, f1 = 0.781, cross_entropy = 12.393
For test set using lr model with seed 18: acc = 0.576, f1 = 0.731, cross_entropy = 14.631
For training set using rf model with seed 18: acc = 0.94, f1 = 0.955, cross_entropy = 2.079

For test set using rf model with seed 18: acc = 0.618, f1 = 0.703, cross_entropy = 13.192
 For training set using lr model with seed 19: acc = 0.637, f1 = 0.778, cross_entropy = 12.553
 For test set using lr model with seed 19: acc = 0.59, f1 = 0.742, cross_entropy = 14.152
 For training set using rf model with seed 19: acc = 0.949, f1 = 0.962, cross_entropy = 1.759
 For test set using rf model with seed 19: acc = 0.59, f1 = 0.697, cross_entropy = 14.152

```
In [94]: fig, axes = plt.subplots(1,2, figsize=(9,5))
ax = axes[0]
ax.boxplot([lr_test_acc, rf_test_acc])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Acc')
ax.set_xlabel('Model')
ax.set_title('Question Answering')

ax = axes[1]
ax.boxplot([lr_test_cross_entropy, rf_test_cross_entropy])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Cross Entropy')
ax.set_xlabel('Model')
ax.set_title('Question Answering')
```

Out[94]: Text(0.5, 1.0, 'Question Answering')



SC

```

In [95]: feature_cat = ['location']
onehot_enc = OneHotEncoder(handle_unknown='ignore')
onehot_enc.fit(data_sc[feature_cat])
sc_cat = onehot_enc.transform(data_sc[feature_cat]).toarray()

feature_num = ['source_label', 'target_label', 'loss', 'test_loss']
scaler = StandardScaler()
scaler.fit(data_sc[feature_num])
sc_num = scaler.transform(data_sc[feature_num])

data_sc['poisoned'] = data_sc['poisoned'].replace({False:0, True:1})

sc_ = pd.concat([pd.DataFrame(sc_cat), pd.DataFrame(sc_num)], axis=1)

features = [
    'start',
    'middle',
    'end',
    'src_label',
    'tgt_label',
    'loss',
    'test_loss']
sc_.columns = features

```

C:\Users\CSY\anaconda3\lib\site-packages\ipykernel_launcher.py:11: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)

This is added back by InteractiveShellApp.init_path()

```

In [96]: lr_test_acc = []
lr_test_cross_entropy = []
rf_test_acc = []
rf_test_cross_entropy = []
for s in range(20):
    X_sc_train, X_sc_test, y_sc_train, y_sc_test = train_test_split(sc_, data_sc_)
    model_lr = LogisticRegression(random_state=1)
    model_lr.fit(X_sc_train, y_sc_train)
    model_rf = RandomForestClassifier(random_state=1)
    model_rf.fit(X_sc_train, y_sc_train)
    score(model_lr, X_sc_train, y_sc_train, 'training', seed=s)
    acc, f1, cross_entropy = score(model_lr, X_sc_test, y_sc_test, 'test', seed=s)
    lr_test_acc.append(acc)
    lr_test_cross_entropy.append(cross_entropy)
    score(model_rf, X_sc_train, y_sc_train, 'training', 'rf', seed=s)
    acc, f1, cross_entropy = score(model_rf, X_sc_test, y_sc_test, 'test', 'rf',
    rf_test_acc.append(acc)
    rf_test_cross_entropy.append(cross_entropy)

```

For training set using lr model with seed 0: acc = 0.778, f1 = 0.714, cross_entropy = 7.675

For test set using lr model with seed 0: acc = 0.8, f1 = 0.719, cross_entropy = 6.908

For training set using rf model with seed 0: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 0: acc = 0.789, f1 = 0.642, cross_entropy = 7.292

For training set using lr model with seed 1: acc = 0.793, f1 = 0.731, cross_entropy = 7.164

For test set using lr model with seed 1: acc = 0.767, f1 = 0.656, cross_entropy = 8.059

For training set using rf model with seed 1: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 1: acc = 0.744, f1 = 0.61, cross_entropy = 8.827

For training set using lr model with seed 2: acc = 0.781, f1 = 0.704, cross_entropy = 7.547

For test set using lr model with seed 2: acc = 0.822, f1 = 0.758, cross_entropy = 6.14

For training set using rf model with seed 2: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 2: acc = 0.733, f1 = 0.625, cross_entropy = 9.21

For training set using lr model with seed 3: acc = 0.8, f1 = 0.716, cross_entropy = 6.908

For test set using lr model with seed 3: acc = 0.811, f1 = 0.738, cross_entropy = 6.524

For training set using rf model with seed 3: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 3: acc = 0.744, f1 = 0.667, cross_entropy = 8.827

For training set using lr model with seed 4: acc = 0.793, f1 = 0.731, cross_entropy = 7.164

For test set using lr model with seed 4: acc = 0.767, f1 = 0.687, cross_entropy = 8.059

For training set using rf model with seed 4: acc = 1.0, f1 = 1.0, cross_entropy

= 0.0
For test set using rf model with seed 4: acc = 0.811, f1 = 0.73, cross_entropy = 6.524
For training set using lr model with seed 5: acc = 0.774, f1 = 0.708, cross_entropy = 7.803
For test set using lr model with seed 5: acc = 0.8, f1 = 0.719, cross_entropy = 6.908
For training set using rf model with seed 5: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 5: acc = 0.756, f1 = 0.676, cross_entropy = 8.443
For training set using lr model with seed 6: acc = 0.789, f1 = 0.716, cross_entropy = 7.292
For test set using lr model with seed 6: acc = 0.811, f1 = 0.761, cross_entropy = 6.524
For training set using rf model with seed 6: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 6: acc = 0.767, f1 = 0.712, cross_entropy = 8.059
For training set using lr model with seed 7: acc = 0.807, f1 = 0.737, cross_entropy = 6.652
For test set using lr model with seed 7: acc = 0.733, f1 = 0.667, cross_entropy = 9.21
For training set using rf model with seed 7: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 7: acc = 0.722, f1 = 0.648, cross_entropy = 9.594
For training set using lr model with seed 8: acc = 0.778, f1 = 0.703, cross_entropy = 7.675
For test set using lr model with seed 8: acc = 0.822, f1 = 0.771, cross_entropy = 6.14
For training set using rf model with seed 8: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 8: acc = 0.733, f1 = 0.667, cross_entropy = 9.21
For training set using lr model with seed 9: acc = 0.811, f1 = 0.727, cross_entropy = 6.524
For test set using lr model with seed 9: acc = 0.767, f1 = 0.712, cross_entropy = 8.059
For training set using rf model with seed 9: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 9: acc = 0.722, f1 = 0.627, cross_entropy = 9.594
For training set using lr model with seed 10: acc = 0.778, f1 = 0.706, cross_entropy = 7.675
For test set using lr model with seed 10: acc = 0.833, f1 = 0.754, cross_entropy = 5.756
For training set using rf model with seed 10: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 10: acc = 0.744, f1 = 0.657, cross_entropy = 8.827
For training set using lr model with seed 11: acc = 0.811, f1 = 0.756, cross_entropy = 6.524
For test set using lr model with seed 11: acc = 0.733, f1 = 0.636, cross_entropy = 9.21
For training set using rf model with seed 11: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 11: acc = 0.767, f1 = 0.696, cross_entropy = 8.059
For training set using lr model with seed 12: acc = 0.815, f1 = 0.752, cross_entropy = 6.396
For test set using lr model with seed 12: acc = 0.722, f1 = 0.648, cross_entropy = 9.594
For training set using rf model with seed 12: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 12: acc = 0.733, f1 = 0.647, cross_entropy = 9.21
For training set using lr model with seed 13: acc = 0.774, f1 = 0.719, cross_entropy = 7.803
For test set using lr model with seed 13: acc = 0.756, f1 = 0.676, cross_entropy = 8.443
For training set using rf model with seed 13: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 13: acc = 0.822, f1 = 0.742, cross_entropy = 6.14
For training set using lr model with seed 14: acc = 0.77, f1 = 0.69, cross_entropy = 7.931
For test set using lr model with seed 14: acc = 0.833, f1 = 0.776, cross_entropy = 5.757
For training set using rf model with seed 14: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 14: acc = 0.767, f1 = 0.677, cross_entropy = 8.059
For training set using lr model with seed 15: acc = 0.793, f1 = 0.741, cross_entropy = 7.164
For test set using lr model with seed 15: acc = 0.733, f1 = 0.625, cross_entropy = 9.21
For training set using rf model with seed 15: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 15: acc = 0.744, f1 = 0.61, cross_entropy = 8.827
For training set using lr model with seed 16: acc = 0.77, f1 = 0.702, cross_entropy = 7.931
For test set using lr model with seed 16: acc = 0.8, f1 = 0.735, cross_entropy = 6.908
For training set using rf model with seed 16: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 16: acc = 0.8, f1 = 0.727, cross_entropy = 6.908
For training set using lr model with seed 17: acc = 0.763, f1 = 0.695, cross_entropy = 8.187
For test set using lr model with seed 17: acc = 0.844, f1 = 0.794, cross_entropy = 5.373
For training set using rf model with seed 17: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 17: acc = 0.756, f1 = 0.686, cross_entropy = 8.443
For training set using lr model with seed 18: acc = 0.815, f1 = 0.731, cross_entropy = 6.396
For test set using lr model with seed 18: acc = 0.789, f1 = 0.765, cross_entropy = 7.292
For training set using rf model with seed 18: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 18: acc = 0.778, f1 = 0.73, cross_entropy

= 7.675

For training set using lr model with seed 19: acc = 0.77, f1 = 0.708, cross_entropy = 7.931

For test set using lr model with seed 19: acc = 0.8, f1 = 0.735, cross_entropy = 6.908

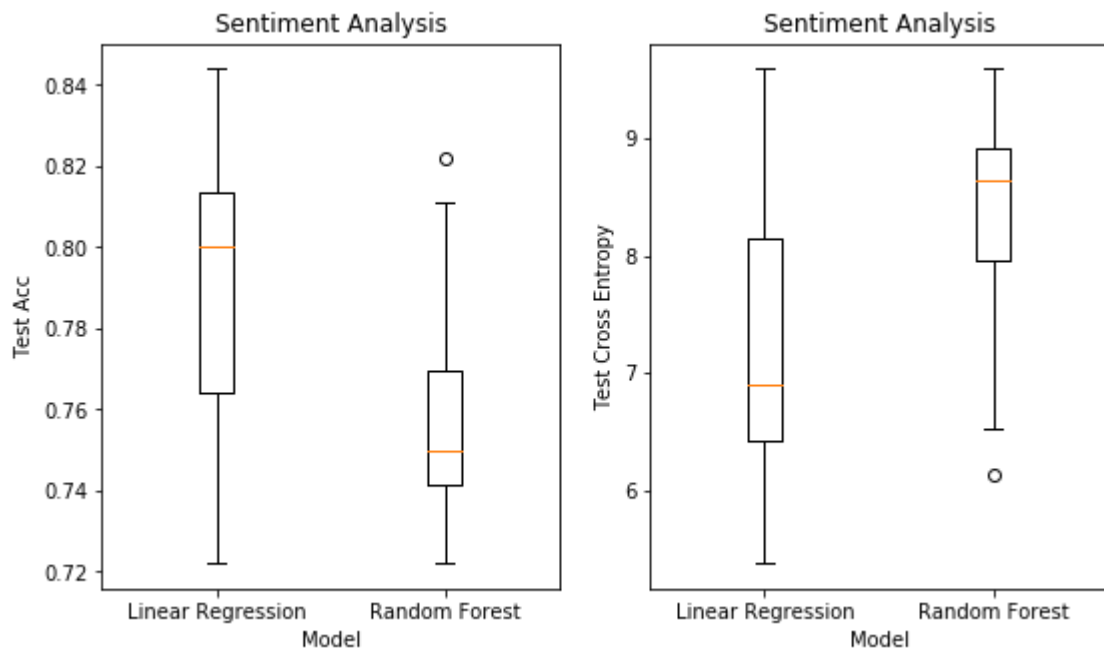
For training set using rf model with seed 19: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 19: acc = 0.744, f1 = 0.623, cross_entropy = 8.827

```
In [97]: fig, axes = plt.subplots(1,2, figsize=(9,5))
ax = axes[0]
ax.boxplot([lr_test_acc, rf_test_acc])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Acc')
ax.set_xlabel('Model')
ax.set_title('Sentiment Analysis')

ax = axes[1]
ax.boxplot([lr_test_cross_entropy, rf_test_cross_entropy])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Cross Entropy')
ax.set_xlabel('Model')
ax.set_title('Sentiment Analysis')
```

Out[97]: Text(0.5, 1.0, 'Sentiment Analysis')



NER

```

In [98]: feature_cat = ['source_label', 'target_label']
onehot_enc = OneHotEncoder(handle_unknown='ignore')
onehot_enc.fit(data_ner[feature_cat])
ner_cat = onehot_enc.transform(data_ner[feature_cat]).toarray()

feature_num = ['loss', 'test_loss']
scaler = StandardScaler()
scaler.fit(data_ner[feature_num])
ner_num = scaler.transform(data_ner[feature_num])

data_ner['poisoned'] = data_ner['poisoned'].replace({False:0, True:1})

ner_ = pd.concat([pd.DataFrame(ner_cat), pd.DataFrame(ner_num)], axis=1)

features = [
    'src_1',
    'src_3',
    'src_5',
    'src_7',
    'tgt_1',
    'tgt_3',
    'tgt_5',
    'tgt_7',
    'loss',
    'test_loss']
ner_.columns = features

```

C:\Users\CSY\anaconda3\lib\site-packages\ipykernel_launcher.py:11: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy> (<http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>)
This is added back by InteractiveShellApp.init_path()

```

In [99]: lr_test_acc = []
lr_test_cross_entropy = []
rf_test_acc = []
rf_test_cross_entropy = []
for s in range(20):
    X_ner_train, X_ner_test, y_ner_train, y_ner_test = train_test_split(ner_, data_ner_,
    model_lr = LogisticRegression(random_state=s)
    model_lr.fit(X_ner_train, y_ner_train)
    model_rf = RandomForestClassifier(random_state=s)
    model_rf.fit(X_ner_train, y_ner_train)
    score(model_lr, X_ner_train, y_ner_train, 'training', seed=s)
    acc, f1, cross_entropy = score(model_lr, X_ner_test, y_ner_test, 'test', seed=s)
    lr_test_acc.append(acc)
    lr_test_cross_entropy.append(cross_entropy)
    score(model_rf, X_ner_train, y_ner_train, 'training', 'rf', seed=s)
    acc, f1, cross_entropy = score(model_rf, X_ner_test, y_ner_test, 'test', 'rf', seed=s)
    rf_test_acc.append(acc)
    rf_test_cross_entropy.append(cross_entropy)

```

For training set using lr model with seed 0: acc = 0.685, f1 = 0.0, cross_entropy = 10.873
 For test set using lr model with seed 0: acc = 0.611, f1 = 0.0, cross_entropy = 13.432
 For training set using rf model with seed 0: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
 For test set using rf model with seed 0: acc = 0.562, f1 = 0.253, cross_entropy = 15.137
 For training set using lr model with seed 1: acc = 0.66, f1 = 0.0, cross_entropy = 11.726
 For test set using lr model with seed 1: acc = 0.685, f1 = 0.0, cross_entropy = 10.873
 For training set using rf model with seed 1: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
 For test set using rf model with seed 1: acc = 0.611, f1 = 0.308, cross_entropy = 13.432
 For training set using lr model with seed 2: acc = 0.679, f1 = 0.0, cross_entropy = 11.087
 For test set using lr model with seed 2: acc = 0.63, f1 = 0.0, cross_entropy = 12.792
 For training set using rf model with seed 2: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
 For test set using rf model with seed 2: acc = 0.586, f1 = 0.23, cross_entropy = 14.285
 For training set using lr model with seed 3: acc = 0.658, f1 = 0.0, cross_entropy = 11.797
 For test set using lr model with seed 3: acc = 0.691, f1 = 0.0, cross_entropy = 10.66
 For training set using rf model with seed 3: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
 For test set using rf model with seed 3: acc = 0.617, f1 = 0.205, cross_entropy = 13.219
 For training set using lr model with seed 4: acc = 0.673, f1 = 0.0, cross_entropy = 11.3
 For test set using lr model with seed 4: acc = 0.648, f1 = 0.0, cross_entropy = 12.153
 For training set using rf model with seed 4: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

opy = 0.0
For test set using rf model with seed 4: acc = 0.617, f1 = 0.295, cross_entropy = 13.219
For training set using lr model with seed 5: acc = 0.685, f1 = 0.0, cross_entropy = 10.873
For test set using lr model with seed 5: acc = 0.611, f1 = 0.0, cross_entropy = 13.432
For training set using rf model with seed 5: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 5: acc = 0.562, f1 = 0.237, cross_entropy = 15.137
For training set using lr model with seed 6: acc = 0.658, f1 = 0.0, cross_entropy = 11.797
For test set using lr model with seed 6: acc = 0.691, f1 = 0.0, cross_entropy = 10.66
For training set using rf model with seed 6: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 6: acc = 0.537, f1 = 0.194, cross_entropy = 15.99
For training set using lr model with seed 7: acc = 0.679, f1 = 0.013, cross_entropy = 11.087
For test set using lr model with seed 7: acc = 0.636, f1 = 0.0, cross_entropy = 12.579
For training set using rf model with seed 7: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 7: acc = 0.574, f1 = 0.289, cross_entropy = 14.711
For training set using lr model with seed 8: acc = 0.681, f1 = 0.0, cross_entropy = 11.015
For test set using lr model with seed 8: acc = 0.623, f1 = 0.0, cross_entropy = 13.005
For training set using rf model with seed 8: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 8: acc = 0.574, f1 = 0.258, cross_entropy = 14.711
For training set using lr model with seed 9: acc = 0.663, f1 = 0.0, cross_entropy = 11.655
For test set using lr model with seed 9: acc = 0.679, f1 = 0.0, cross_entropy = 11.087
For training set using rf model with seed 9: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 9: acc = 0.574, f1 = 0.127, cross_entropy = 14.711
For training set using lr model with seed 10: acc = 0.679, f1 = 0.0, cross_entropy = 11.087
For test set using lr model with seed 10: acc = 0.63, f1 = 0.0, cross_entropy = 12.792
For training set using rf model with seed 10: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 10: acc = 0.531, f1 = 0.174, cross_entropy = 16.203
For training set using lr model with seed 11: acc = 0.665, f1 = 0.0, cross_entropy = 11.584
For test set using lr model with seed 11: acc = 0.673, f1 = 0.0, cross_entropy = 11.3
For training set using rf model with seed 11: acc = 1.0, f1 = 1.0, cross_entropy = 0.0

For test set using rf model with seed 11: acc = 0.623, f1 = 0.265, cross_entropy = 13.005
For training set using lr model with seed 12: acc = 0.679, f1 = 0.0, cross_entropy = 11.087
For test set using lr model with seed 12: acc = 0.63, f1 = 0.0, cross_entropy = 12.792
For training set using rf model with seed 12: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 12: acc = 0.58, f1 = 0.244, cross_entropy = 14.498
For training set using lr model with seed 13: acc = 0.677, f1 = 0.0, cross_entropy = 11.158
For test set using lr model with seed 13: acc = 0.636, f1 = 0.0, cross_entropy = 12.579
For training set using rf model with seed 13: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 13: acc = 0.623, f1 = 0.265, cross_entropy = 13.005
For training set using lr model with seed 14: acc = 0.658, f1 = 0.0, cross_entropy = 11.797
For test set using lr model with seed 14: acc = 0.691, f1 = 0.0, cross_entropy = 10.66
For training set using rf model with seed 14: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 14: acc = 0.58, f1 = 0.333, cross_entropy = 14.498
For training set using lr model with seed 15: acc = 0.648, f1 = 0.0, cross_entropy = 12.153
For test set using lr model with seed 15: acc = 0.722, f1 = 0.0, cross_entropy = 9.594
For training set using rf model with seed 15: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 15: acc = 0.722, f1 = 0.43, cross_entropy = 9.594
For training set using lr model with seed 16: acc = 0.652, f1 = 0.0, cross_entropy = 12.01
For test set using lr model with seed 16: acc = 0.71, f1 = 0.0, cross_entropy = 10.021
For training set using rf model with seed 16: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 16: acc = 0.642, f1 = 0.356, cross_entropy = 12.366
For training set using lr model with seed 17: acc = 0.681, f1 = 0.0, cross_entropy = 11.015
For test set using lr model with seed 17: acc = 0.623, f1 = 0.0, cross_entropy = 13.005
For training set using rf model with seed 17: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 17: acc = 0.586, f1 = 0.247, cross_entropy = 14.285
For training set using lr model with seed 18: acc = 0.656, f1 = 0.012, cross_entropy = 11.868
For test set using lr model with seed 18: acc = 0.704, f1 = 0.0, cross_entropy = 10.234
For training set using rf model with seed 18: acc = 1.0, f1 = 1.0, cross_entropy = 0.0
For test set using rf model with seed 18: acc = 0.623, f1 = 0.299, cross_entropy = 13.005

```

ropy = 13.005
For training set using lr model with seed 19: acc = 0.681, f1 = 0.0, cross_e
ntropy = 11.015
For test set using lr model with seed 19: acc = 0.611, f1 = 0.0, cross_entro
py = 13.432
For training set using rf model with seed 19: acc = 1.0, f1 = 1.0, cross_ent
ropy = 0.0
For test set using rf model with seed 19: acc = 0.556, f1 = 0.217, cross_ent
ropy = 15.351

```

```

In [100]: fig, axes = plt.subplots(1,2, figsize=(9,5))
ax = axes[0]
ax.boxplot([lr_test_acc, rf_test_acc])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Acc')
ax.set_xlabel('Model')
ax.set_title('Name Entity Recognition')

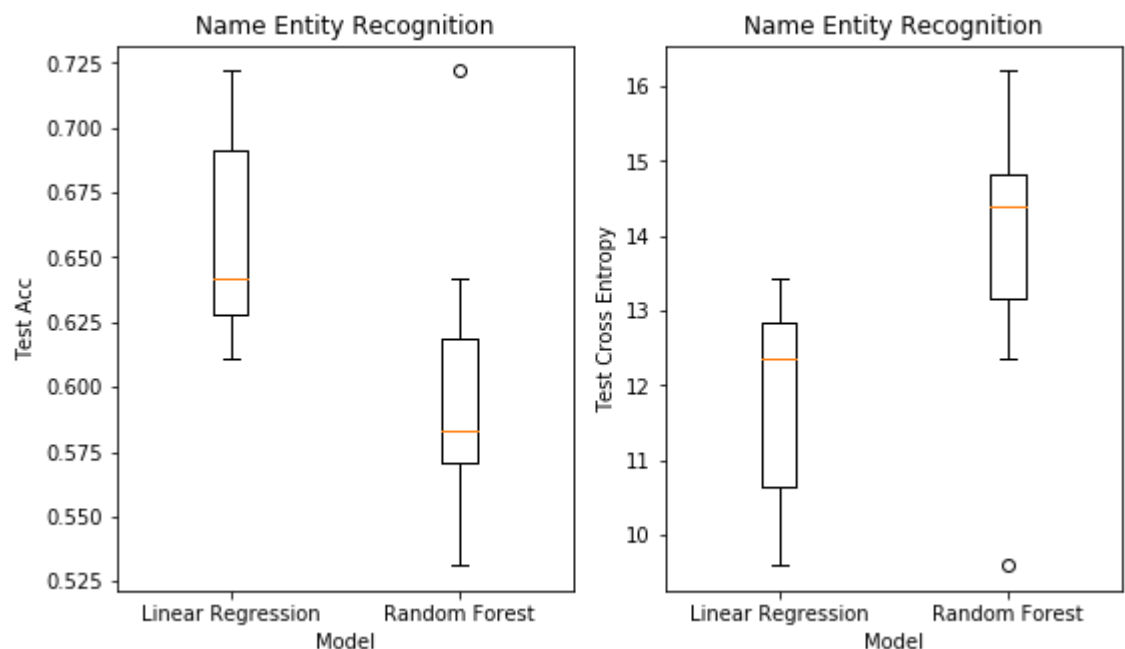
ax = axes[1]
ax.boxplot([lr_test_cross_entropy, rf_test_cross_entropy])
ax.set_xticklabels(['Linear Regression', 'Random Forest'])
ax.set_ylabel('Test Cross Entropy')
ax.set_xlabel('Model')
ax.set_title('Name Entity Recognition')

```

```

Out[100]: Text(0.5, 1.0, 'Name Entity Recognition')

```



```

In [ ]:

```

